

How Lmod Loads a Modulefile, Part 1

Robert McLay

June. 7, 2022

Outline



- ▶ Talk: How Lmod finds a modulefile to load
- ▶ Next Time: Part 2: How Lmod actually evaluates a modulefile
- ▶ There is a surprising amount to talk about before the modulefile gets evaluated.
- ▶ We will be talking about a user requested load.
- ▶ A module requested load is slightly different.

Outline (II)

- ▶ `main()`: How `main()` parses the command line
- ▶ `Load_Usr()`: Starts working your modules to load
- ▶ `l_usrLoad()`: Splits modules to loads or unloads
- ▶ `MName Class`: Maps module name to a filename
- ▶ `mcp`: Object to control kind of evaluation
- ▶ `MasterControl:load(mA)`: We are loading a modulefile
- ▶ `Master:load(mA)`: Where the heavy lifting is done
- ▶ `loadModuleFile()`: Where the modulefile is actually evaluated

main(): parsing command line

- ▶ main() is located in lmod.in.lua (installed as lmod)
- ▶ The command is: *module load foo*
- ▶ To parse Lmod uses the 2nd argument w/o a minus: *load*
- ▶ Lmod searches lmodCmdA to map string to command

A side note on Lmod Coding Conventions

- ▶ A variable w/ a trailing “A” means that an array
- ▶ A variable w/ a trailing “T” means that an table (or dictionary)
- ▶ A variable w/ a trailing “Tbl” means that an table (or dictionary)
- ▶ A routine with a l_name is a local function (file scope)
- ▶ Class Name are in CamelCase

lmodCmdA

```
local lmodCmdA = {  
  {cmd = 'add',          min = 2, action = loadTbl      },  
  {cmd = 'avail',        min = 2, action = availTbl      },  
  {cmd = 'isloaded',      min = 3, action = isLoadedTbl  },  
  {cmd = 'is_loaded',     min = 4, action = isLoadedTbl  },  
  {cmd = 'is-loaded',     min = 4, action = isLoadedTbl  },  
  {cmd = 'load',          min = 2, action = loadTbl      },  
  ...  
}  
local loadTbl      = { name = "load", cmd = Load_Usr    }
```

- So “load” matches “load” with more than 2 chars

Going from lmodCmdA to Load_Usr()

- ▶ Sets cmdT to loadTbl
- ▶ Sets cmdName to cmdT.name (forces standard name not user command name)
- ▶ cmdT.cmd(unpack(masterTbl.pargs)) \Rightarrow Load_Usr

Calling Load_Usr()

- All functions implementing user commands are in `src/cmdfunc.lua`

```
function Load_Try(...)
    dbg.start"Load_Try(",concatTbl({...}," ",")"
    local check_must_load = false
    local argA            = pack(...)
    l_usrLoad(argA, check_must_load)
    dbg.fini("Load_Try")
end

function Load_Usr(...)
    dbg.start"Load_Usr(",concatTbl({...}," ",")"
    local check_must_load = true
    local argA            = pack(...)
    l_usrLoad(argA, check_must_load)
    dbg.fini("Load_Usr")
end
```


l_usrLoad(argA, check_must_load)

- ▶ Split argA into loads in lA, unloads in uA (-foo)
- ▶ Both uA and lA are an array of MName objects.
- ▶ unload modules in uA
- ▶ `lA[#lA+1] = MName:new("load", module_name)`
- ▶ `mcp:load_usr(lA)`
- ▶ `src: src/cmdfunc.lua`

MName class: Module Name class

- ▶ Maps name (“foo” or “foo/1.1”) to filename
- ▶ There are two kinds of searching “load” or “mt”
- ▶ Load: must search file system.
- ▶ mt: filename is in modultable
- ▶ Evaluation must be lazy or just-in-time
- ▶ Software hierarchy means that
- ▶ `module load gcc mpich`
- ▶ mpich might not be in \$MODULEPATH until after gcc is loaded
- ▶ src: src/MName.lua

MName key concepts

- ▶ **userName:** name on the command line
- ▶ It might be gcc or gcc/9.3.0
- ▶ **sn:** the shortName or a name without a version
- ▶ **fullName:** The full name of the module (sn/version)
- ▶ Examples:
 1. gcc/9.3.0 (sn: gcc, N/V)
 2. gcc/x86_64/9.3.0 (sn: gcc, N/V/V)
 3. compiler/gcc/9.3.0 (sn: compiler/gcc, C/N/V)
 4. compiler/gcc/x86_64/9.3.0 (sn: compiler/gcc, N/V/V)

mcp and MasterControl class

- ▶ MasterControl class is what controls whether a "load" in a modulefile is a load or unload
- ▶ mcp is a global variable that is built to be in a mode() like load, unload, spider, etc.
- ▶ We talked about this in an earlier presentation.

MasterControl:load(mA)

```
function M.load(self, mA)
  local master = Master:singleton()
  local a      = master:load(mA)

  if (not quiet()) then
    self:registerAdminMsg(mA) -- nag msg registration.
  end
  return a
end
```

- ▶ MasterControl functions call Master Functions to do the work.
- ▶ src: src/MasterControl.lua

Master:load(mA)

```
function M.load(mA)
  for i = 1, #mA do
    repeat
      mname = mA[i]
      sn = mname:sn() -- shortName
      fn = mname:fn() -- file name
      -- if blank sn -> pushModule (might have to wait until
      -- compiler or mpi is loaded.
      -- and break (really continue)

      if (mt:have(sn,"active")) then
        -- Block 1: Check for previously loaded module with same sn

      elseif (fn) then
        -- Block 2: Load modulefile

        -- Check for family stack (e.g. compiler, mpi etc)
        if (mcp.processFamilyStack(fullName)) then
          -- Suppose gcc is loaded and it was "replaced" by intel
          -- unload gcc and reload intel
        end
      until true
    end
  end

  -- Reload every module if change in MODULEPATH.

  -- load any modules on module stack
end
```

► This is where the heavy lifting is done.

Block 1: Check for previous loaded module w/same sn

```
if (mt:have(sn,"active")) then
  -- if disable_same_name_autoswap -> error out
  -- Otherwise: unload previous module
  local mcp_old = mcp
  local mcp      = MCP
  unload_internalMName:new("mt",sn)
  mname:reset() -- force a new lazyEval
  local status = mcp:load_usr{mname}
  mcp           = mcp_old
```

- ▶ Here we guarantee the right mcp
- ▶ Unload the old module
- ▶ Recursively mcp:load_usrmname

Block 2: Load modulefile

```
elseif (fn) then
  frameStk:push(mname)
  mt = frameStk:mt()
  mt:add(mname,"pending")
  local status = loadModuleFile{file = fn, shell = shellNm,
                                mList = mList, reportErr = true}
  frameStk:pop()
  loaded = true
end
```


loadModuleFile(t)

- ▶ This is where Lmod handle either *.lua files or TCL Modulefiles
- ▶ Once either read in as a block (for *.lua) or converted (TCL modulefile \Rightarrow Lua)
- ▶ src: src/loadModulefile.lua

```
-- Use the sandbox to evaluate modulefile text.  
if (whole) then  
    status, msg = sandbox_run(whole)  
else  
    status = nil  
    msg     = "Empty or non-existent file"  
end  
-- report any errors
```

Next time

- ▶ What is a sandbox and how does it work?
- ▶ Why I want a sandbox?
- ▶ Next time handing control to modulefile

Conclusions

- ▶ It takes a lot to get to the point where Lmod is evaluating your modulefile.
- ▶ Lmod uses several “classes” to manage the loading of a module
- ▶ Plus a couple of Design Patterns such as Singletons

Future Topics

- ▶ Next Meeting: July 5th 9:30 US Central (14:30 UTC)
- ▶ What happens from the `loadModulefile(t)`.
- ▶ This is where Lmod hands off control to the user.